

- A dict mapping i/p names to the corresponding array/tensors.
- A tf.data dataset. Should return a tuple of either (inputs, targets) or (iips, target, sample-weight).
- A generator or keras.utils.Sequence returning the above.

$Y \rightarrow$ ◦ Needs to be in the same format as 'X'.

◦ We should shuffle the data to get rid of any imposed order.

\rightarrow from sklearn.utils import shuffle

train_labels, train_samples = shuffle (train_labels, train_samples)

◦ Also for scaling b/w 0 to 1:-

\rightarrow from sklearn.preprocessing import MinMaxScaler

scaler = MinMaxScaler(feature_range=(0,1))

↑ — MinMaxScaler Object

Scaled_train_samples = scaler.fit_transform(train_samples.reshape(-1, 1))

→ reshape (-1, 1) returns a
2D array w/ 1 col,
and as many rows necessary
to accomodate the data.

The MinMaxScaler doesn't take
1D array by default, so we
pass it through this to transform
it.

** Calling .shape on a numpy
array returns its shape.

Ex:- Array w/ 100 rows, & 1 col. \Rightarrow 2 rows, 1 cols :- [[1], [2]]
 \Rightarrow (2, 1)
 \Rightarrow arr.shape
 $\Rightarrow (100, 1)$

Building the model :-

```
import tensorflow as tf
from tensorflow import keras
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Activation, Dense
from tensorflow.keras.optimizers import Adam
```

```
model = Sequential([
    ① Dense(units=16, input_shape=(1,), activation='relu'),
    ② Dense(units=32, activation='relu'),
    ③ Dense(units=2, activation='softmax')
])
```

I

16 nodes (arbitrary) → relu follows this layer
① 16 → arb.
② 32 → arb.
③ 2 → for percentage (prob.)

The ① Dense layer is actually not the first layer.
→ It is the first hidden layer.

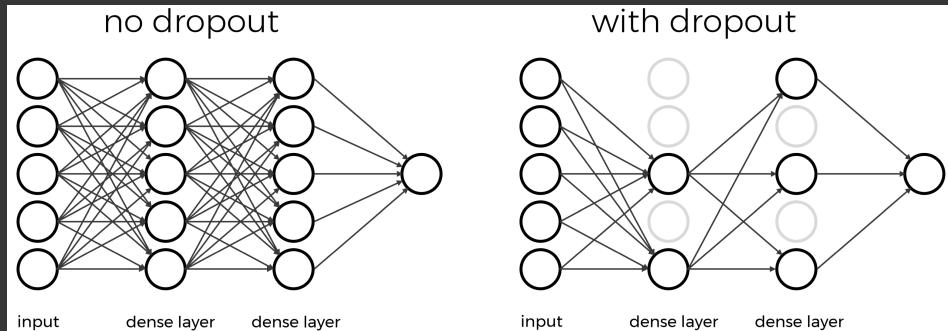
→ We're not explicitly defining the ip layer,
it is specified by the input_shape parameter.

```
model.summary()
Model: "sequential"
Layer (type)          Output Shape         Param #
dense (Dense)        (None, 16)           32
dense_1 (Dense)       (None, 32)           544
dense_2 (Dense)       (None, 2)            66
Total params: 642
Trainable params: 642
Non-trainable params: 0
```

$$\begin{aligned} & 16 + 16 \text{ bias} \\ & 16 * 32 + 32 \text{ bias} \\ & 32 * 2 + 2 \text{ bias} \end{aligned}$$

importing the Sequential model from sequential api | Contrast this to the functional api

Dense Layer



→ Dense Layer refers to a fully connected layer.

→ It may have dropout
→ Called so b/c they're densely connected

> model.compile(optimizer=Adam(learning_rate=0.001), loss='Sparse categorical Crossentropy',
 metrics=['accuracy'])

Common way to evaluate performance

① → optimizers are used to minimize the loss fcn

→ Adam:- Adaptive Momentum Estimation

↳ can be thought of as a combination of RMS prop and Stochastic Gradient Descent

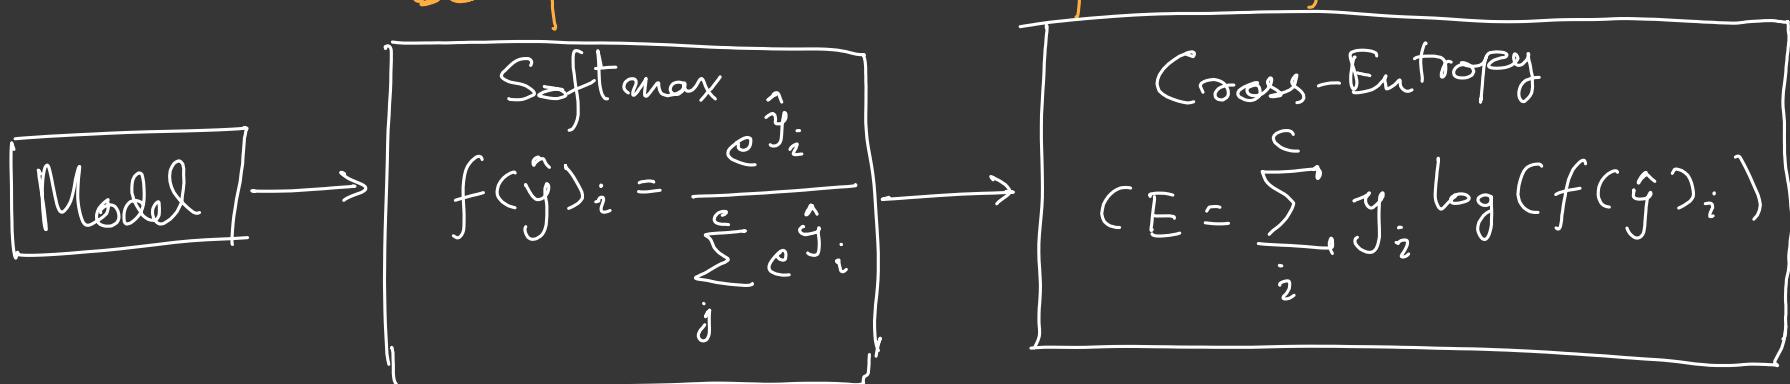
② Sparse Categorical Cross-Entropy

→ Extension of Categorical Cross Entropy, where o/p labels are represented in Sparse matrix format.

◦ Categorical Cross-Entropy:- Softmax loss

→ Softmax Activation + Cross Entropy loss

→ Used for Multiclass Classification. (for CNNs)



```
> model.fit(x=scaled_train_samples, y=train_labels, validation_split=0.1,  
batch_size=10, epochs=30, shuffle=True, verbose=2)
```

$x \rightarrow$ already done

$y \rightarrow$ target_data, done

batch size \rightarrow how many data points to be included
in one pass

epochs \rightarrow ✓

shuffle \rightarrow By default already True,
make False for Time Series

verbose = [0, 1, 2] $\xrightarrow{\text{most verbose}}$
 \downarrow
no verbose

Epoch 1/30

2100 / 2100

\uparrow this number is total data
divided by batch
size

Validation Set

- The model is still learning on or 'training on' the training data
- It allows us to see how general our model is/ how well it is able to generalize on data. This is done for each epoch.

* → So, if we see the model is giving very good result for the training set, but sub-par results for the validation set, then the model is overfit.

- There are 2 ways of creating a validation set:-

model.fit (x=scaled_train_samples, y= train_labels,

① Validation_data= valid_set, ② Validation_split = 0.1,

batch_size = 10,

epochs = 30,

* shuffle = True,
verbose = 2)

→ The shuffle function is called after splitting the data, so let's say we have a dataset where it is arranged by age, so the last 10% will have one parameter not shuffled throughout, so it is important to pre-shuffle the data before calling split_validation.

Testing - (Inference)

- Make a test set.
- Make it into a numpy array.
- Then shuffle it
- Then reshape it and scale it accordingly

Predict:-

`predictions = model.predict(x=scaled_test_samples, batch_size=10, verbose=0)`

→ This gives us a number b/w 0 & 1

returns an array [$\xrightarrow{\text{no side eff.}}$ 0.92, $\xrightarrow{\text{side eff.}}$ 0.077]

→ should be the same
batch size

→ Take the index of the higher index for only the prediction w/ the higher prob.

→ $\begin{matrix} 0 \\ \vdots \\ 1 \dots \end{matrix}$

Confusion Matrix

→ It is in sklearn

%matplotlib inline

```
from sklearn.metrics import confusion_matrix
```

```
import tensorflow
```

```
import matplotlib.pyplot as plt
```

```
cm = confusion_matrix(y_true = test_labels, y_pred = rounded_predictions)
```

* → def plot_confusion_matrix(...)

on scikit learn

website

```
cm_plot_labels = ['no-side-effects', 'had-side-effects']
```

```
plot_confusion_matrix(cm=cm, classes=cm_plot_labels, title='Confusion Matrix')
```

Saving & Loading a Model

Saving:-

```
import os, path
```

```
if os.path.isfile('models/model_name.h5') is False:  
    model.save('models/model_name.h5')
```

This save function saves -

- The architecture of the model
- The weights of the model
- The training configuration
- The state of the optimizer , allowing to resume training exactly where it was left off.

Loading :-

```
from tensorflow.keras.models import load_model
```

```
new_model = load_model('models/name.h5')
```

`model.get_weights()` → gives the weights of the model

From & To JSON or YAML

`json_string = model.to_json()`

`yaml_string = model.to_yaml()`

`from tensorflow.keras.models import model_from_json`

`model_arch = model_from_json(json_string)`

→ This does not save the weights for the model.

→ We need `model.save_weights()` for that.

`if os.path.isfile('models/model_name.h5') is False:`

`model.save_weights('models/model_name.h5')`

Convolutional Neural Networks

